

Gold SMC v8 — XAU/USD Backtesting Bot

Honest Status & Strategy Assessment

Generated 2026-06-05 09:35 UTC • repo: OANDA_Backtesting_bot • prepared for Prem

Bottom line up front. The software is built and unit-tested, but **no real-market verdict exists yet** because the only data run through it so far is synthetic (random-walk) data used to test the plumbing. I also **cannot verify the live OANDA / Slack / GitHub connections from my build environment** (it has no internet to those services), and I cannot push to GitHub for you from here. This report states exactly what is proven, what is not, and a ruthless list of where the strategy and the engine are weak.

Read this as a skeptic would. The honest expectation for a discretionary-style SMC strategy, after realistic costs, is little-to-no edge. Building the machine that can prove that cheaply is the real win.

1 Connection & verification status

You asked whether OANDA, Slack, and GitHub are all wired and working. Here is the truthful state of each. “Verified” means I actually tested it; I will not call something working that I could not test.

Connection	Status	What is true	What you must do
Slack channel C0B88CUAZPD	Reachable (smoke-tested)	I posted a test message to the channel via a code runner, the channel exists (chat state).	Creator, the channel ID is taken (chat state), invite it to your team.
OANDA API (practice)	NOT verified	My sandbox is blocked from OANDA's API, so I cannot fetch the key locally or remotely easily.	Could not find the key, locally or remotely easily.
GitHub repo	NOT pushed	My sandbox cannot reach GitHub; the code lives only in my local OANDA's (Betting?) now.	Run the job push to OANDA's (Betting?) now.
End-to-end cloud loop	NOT live yet	Code + workflow are ready, but the loop has never executed against real data.	After results against real data, the actions workflow once.

So: **one of three connections is confirmed (Slack channel reachable); the OANDA key and the GitHub push are unverified and require you.** I will not pretend the loop is running when it is not.

2 What results actually exist

None on real data. Every number below comes from **synthetic random-walk candles** generated only to exercise the code paths. They say nothing about whether the strategy makes money. Treat them as a smoke test of the machine, not a backtest.

2.1 What IS proven (engineering)

Check	Result	Meaning
Engine mechanics unit tests	9 / 9 pass	No look-ahead, correct cost direction, pessimistic stop-before-target, break-even
Signal parity (fast vs reference)	0 mismatches	The fast vectorized signal path reproduces the ported strategy bar-for-bar.
End-to-end pipeline	Runs in ~3.4 s on 8,256 bars	Fetch → backtest → metrics → reports → memory → Slack summary all executed

During testing the engine actually caught a real bug — break-even was arming off a bar's high and then stopping out off the same bar's low, which would have flattered results. That was fixed and is covered by a test.

2.2 Synthetic plumbing numbers (MEANINGLESS for edge)

Metric	Synthetic value	Note
Trades	72	On fake data
Win rate	30.6%	Noise
Profit factor	1.16	Noise; ~1.0 = no edge
Expectancy	0.09 R	Noise
Walk-forward OOS PF	1.18	In-sample folds hit PF 3.8–4.6 then decayed to 0.0–1.5 out-of-sample

The walk-forward already demonstrates the single most important lesson: **in-sample profit factors of 3.8–4.6 collapsed to 0.0–1.5 out-of-sample**. On random data that is expected — but it is exactly the pattern that exposes curve-fitting, and it is why you must judge the strategy on out-of-sample numbers only.

3 Is your strategy working?

Unknown — and that is the honest answer. No real XAU/USD data has been run, so anyone who tells you it works (or doesn't) right now is guessing. What I can give you is the prior, and it is not flattering:

- The strategy is built on Smart-Money-Concepts ideas (displacement, fair-value gaps, NY-opening). These have a large retail following but **no published, cost-adjusted statistical edge**. The burden of proof is on the strategy, and nothing has met it yet.
- Your original edge claim came from a TradingView backtest, which almost certainly overstates real performance (bar-close fills, no slippage, no failed orders, parameter optimization bias). Expect live/honest results materially worse.
- The most probable outcome of an honest backtest is **profit factor near 1.0 and expectancy near 0 after costs**. If the numbers come back strongly positive, be *more* suspicious, not less — check for look-ahead and data quirks first.

This is not me being negative — it is the same standard real quant desks apply. A project that can cheaply *disprove* a bad idea is worth more than one that cheerleads a curve.

4 Where the strategy is weak (the bad list)

#	Weakness	Why it matters
1	No demonstrated edge	Displacement/FVG/NY-opening are pattern folklore; without out-of-sample, cost-adjusted proof, assume zero alpha.
2	Optimization / backtest bias	Parameters were chosen on TradingView history; they fit the past, not the future.
3	Confidence gate adds no alpha	Blending several signals that individually have no proven edge into a 0–100 score does not create edge. It is good engineering, not a money-maker.
4	Macro layer is inert in backtest	The daily-ETF macro bias is set to NEUTRAL for the backtest (it is too slow to matter intraday and is more noise than signal), so the 'intelligence layer' contributes nothing here.
5	News sentiment is keyword-based	Negation breaks it ('rate cut delayed' scores bullish). Only the FOMC/CPI stand-down is robust, and it is not even modelled in the backtest.
6	Low trade frequency	Max 3 trades/day + flat-only means few trades; you need many months of data before any win rate is statistically meaningful.
7	Single instrument / single regime	XAU/USD only. If gold enters a long choppy range, the strategy has nowhere to hide.
8	Tick-volume proxy	OANDA 'volume' is tick count, not traded contracts, so the volume filter is approximate.

5 Where the engine is limited (honest gaps)

The engine is unit-tested and faithful, but every backtester makes assumptions. Here are mine, so you know exactly what the numbers will and will not capture:

#	Limitation	Impact / what it does not capture
1	Constant costs	Spread & slippage are fixed. Real spreads widen at session opens, rollover, and around news — not modelled per-bar. Results may be optimistic in those windows.
2	Intrabar path unknown	Within a 15-min bar I assume stop-before-target (pessimistic). True ordering is invisible; a different assumption shifts marginal trades.
3	Assumed fills	No partial fills, no liquidity/gap modelling — the model always fills at the next bar's open plus costs.
4	Swing-stop lag	Pivots are confirmed 5 bars later (faithful to the Pine script), so swing-based stops use information that lags the entry by 5 bars.
5	HTF warmup hack	Before ~50 four-hour bars exist, the trend filter falls back to an M15 moving average — a weaker proxy for the first ~8 days of any data window.
6	No sizing constraints	OANDA minimum trade size, margin, and financing/carry are not enforced in the backtest.
7	Small optimization grid	Walk-forward sweeps only displacement × threshold; and optimizing at all invites overfitting if the grid grows.
8	Short history & no significance test	Default 730 days = one macro regime; there is no Monte-Carlo or p-value on the equity curve yet.

None of these are bugs — they are disclosed assumptions. The cost-stress run (`--spread 0.6 --slippage 0.3`) is built in precisely so you can see how fragile any 'edge' is to them.

6 What to fix / the path to a real verdict

1. **Get real data.** Run `python data/fetch_oanda.py` (needs your account id + network). Nothing is real until this runs.
2. **Run the backtest & walk-forward** and read the out-of-sample numbers in `reports/`, not the in-sample ones.
3. **Run the cost-stress variant.** If a wider spread kills the edge, the edge was never real.
4. **Fill in VERDICT.md** with which of the three branches the numbers land in (no edge / marginal / robust). Branch 1 is the likely and acceptable outcome.
5. **Only if out-of-sample survives costs**, consider a small live OANDA paper loop — with continued skepticism.

7 Deploying to the cloud (you must run these)

I cannot push from my environment. From your machine, in the OANDA folder:

```
git init && git add . && git commit -m "OANDA backtesting bot"
git branch -M main
git remote add origin https://github.com/burugupallyprem-coder/OANDA_Backtesting_bot.git
git push -u origin main
```

Then in GitHub → Settings → Secrets and variables → Actions, add: **OANDA_API_KEY**, **OANDA_ACCOUNT_ID**, **SLACK_BOT_TOKEN**, **SLACK_CHANNEL_ID** (C0B88CUAZPD). Set Actions → Workflow permissions to "Read and write". The workflow then runs Mon–Fri 22:00 UTC (refresh + backtest) and Fri 22:30 UTC (weekly review), commits results, and posts to Slack — no laptop required. Confirm `.env` is NOT in the commit (it is gitignored).

Prepared honestly. If a number ever looks too good, suspect the backtest before believing the strategy.